

What if income shifts from labor to capital?

The fiscal and distributional incidence of AI-driven growth in the United States*

Max Ghenis[†]

July 2026

Abstract

I study how AI-driven growth would move United States government revenue, benefit outlays, poverty and inequality, by running two experiment families through PolicyEngine-US, an open-source static microsimulation of the full federal and state tax and transfer system, on openly published microdata. A mechanism sweep shifts 0–100% of labor income into capital income under current law; a scenario set adopts the Budget Lab’s calibration of the [Karger et al. \(2026\)](#) forecaster survey — Slow, Moderate and Rapid adoption crossed with three wage-inequality variants — so that remaining differences between the two models are differences in coverage, not assumptions. In Rapid, the forecast tilt toward capital keeps 36% of the revenue gain that shares-fixed growth would deliver (+211B against +581B) and erases the poverty gain from growth (-1.09pp becomes +0.11pp). Across the wage-inequality variants, net revenue stays within +179B–+244B while the poverty outcome flips sign, from -1.54pp to +1.72pp; automatic transfer responses absorb 47% of the gross-tax swing across those variants, \$42B of it in programs outside a tax-only model, and state tax doubles from +24B to +48B. If less than 22.5% of the incremental capital income is realized on tax returns, the household sector loses revenue on Rapid AI growth. Reconciled cell by cell against the Budget Lab’s published grid, payroll responses agree closely throughout, while the bridged headline is +301B against their +216B (1.40×), concentrated in the income-tax response to the capital shock.

Keywords: artificial intelligence, factor shares, microsimulation, taxation, transfers, poverty, inequality, state taxes

JEL codes: H24, H53, H71, D31, I32, O33, C63

***Acknowledgements:** the cell-level reconciliation in Section 6 is possible because the Budget Lab at Yale publishes its code and complete result grids alongside its reports; I thank them for that practice. All errors are mine. An interactive companion to this paper is at policyengine.org/us/ai-inequality/income-shift; code and data references are in Appendix C.

[†]PolicyEngine. Email: max@policyengine.org.

1 Introduction

Forecasters expect artificial intelligence to raise output and, at the same time, to tilt its composition away from labor and toward capital (Karger et al., 2026). The United States taxes those two kinds of income very differently, and it runs a transfer system that responds automatically when wages fall. So the fiscal and household consequences of AI-driven growth depend not mainly on how much growth arrives, but on which incomes it arrives in — and on which parts of the fiscal system a model can see. The Budget Lab at Yale recently scored AI growth scenarios through a federal individual income and payroll tax microsimulation and found that a growing capital share mutes the revenue gains (Iselin and Nunn, 2026). Their model, by design, stops at the tax code: no transfer programs beyond tax credits, no poverty measurement, no state systems.

This paper runs the same question through the whole system. PolicyEngine-US is an open-source static microsimulation of United States federal and state tax and benefit law — income and payroll taxes, refundable credits, SNAP, SSI, TANF, WIC and the other major transfer programs, and the Supplemental Poverty Measure (SPM) — running on openly published microdata. I use it in two complementary designs. The first is a mechanism sweep: shift 0–100% of positive labor income into capital income, holding modelled market income fixed, and record what current law does at each step. The second is a calibrated scenario set that deliberately adopts the Budget Lab’s implementation of the Karger et al. (2026) forecaster survey — Slow, Moderate and Rapid adoption paths crossed with compressive, proportional and expansive wage-inequality variants, plus their shares-fixed counterfactuals — so that where the two models differ, the difference is model coverage rather than scenario assumptions. Because they publish their code and their complete result grid, the comparison can then be made cell by cell on an identical accounting basis.

Four findings organize the paper. First, the composition tilt, not the growth rate, drives the fiscal and poverty results. Rapid adoption with factor shares held at today’s levels raises net government revenue by +581B in 2030 and cuts SPM poverty by -1.09pp; the same growth with the forecast tilt keeps 36% of the revenue (+211B) and erases the poverty gain (+0.11pp).

Second, the assumption the forecasting literature is least sure of — whether AI compresses or widens the wage distribution — decides the household verdict while leaving the revenue verdict intact. Across the Budget Lab’s own compressive-to-expansive span, Rapid net revenue stays between +179B and +244B, always positive, while the poverty outcome flips sign: -1.54pp under compression, +1.72pp under expansion, a 3.26pp span.

Third, the parts of the fiscal system outside a tax-only model respond materially, and in both directions. Across those same wage-spread variants, gross tax collections swing \$122B, and automatic transfer responses absorb 47% of that swing — \$42B of it in non-credit programs such as SNAP and SSI that a tax microsimulation does not carry. State tax codes add +36B under proportional Rapid growth and double across the variants, from +24B to +48B. Whether these channels matter is exactly the kind of question a full-system model exists to answer.

Fourth, two definitional choices about the tax base move the revenue answer by more than the choice between Slow, Moderate and Rapid. If AI’s incremental capital income goes unrealized — retained, accrued, or sheltered inside retirement accounts — revenue falls roughly linearly, and below a 22.5% realized share the household sector loses money on Rapid AI growth; poverty, meanwhile, sits still to within 12.27–12.24 pp across the whole sweep. And whether taxable retirement distributions count as capital income halves the Rapid estimate when excluded (+211B to +106B).

The reconciliation with the Budget Lab in Section 6 is, to my knowledge, the first cell-level comparison of two independent microsimulations of AI fiscal scenarios. It is reassuring where the models share machinery — payroll responses agree closely in every cell, two independent wage distributions meeting the same taxable maximum — and informative where they do not: the bridged headline is +301B against their +216B ($1.40\times$), with the gap concentrated in the income-tax response to the capital shock ($1.97\times$ on the cell where the wage-spread channel is neutral), consistent with a larger realized-capital base and a more top-concentrated allocation rule.

The paper proceeds as follows. Section 2 places the exercise in the literature. Section 3 describes the model, data, and both experiment designs. Section 4 reports the sweep and scenario results, including the transfer-system and state channels. Section 5 reports the realization and capital-scope sensitivities and stability across data builds. Section 6 reconciles with the Budget Lab cell by cell. Section 7 discusses limitations and next steps.

2 Related work

The paper sits at the intersection of three literatures: macro-fiscal analyses of AI, tax-benefit microsimulation of technology shocks, and the factor-share literature that motivates the experiment’s central mechanism.

AI and the public finances. Iselin and Nunn (2026) is the direct antecedent: they translate the Karger et al. (2026) forecaster survey into 2030 growth, factor-share and wage-distribution scenarios and score them through a federal individual income and payroll tax microsimulation built on the IRS Public Use File, adding a corporate revenue wedge outside the microsimulation. Their central finding — that growth accruing to capital raises less revenue than the same growth at fixed factor shares, by roughly a factor of two in the Rapid scenario — is the claim this paper re-derives on independent machinery and extends to transfers, poverty and states. Karger et al. (2026) elicit forecasts from economists, AI-company staff, policy researchers, superforecasters and the public; the median expects modest GDP acceleration by 2030 with substantial disagreement about wage distribution effects — disagreement this paper propagates rather than resolves. Acemoglu (2025) argues on task-level grounds for modest aggregate effects; Bastani and Waldenström (2024) survey the tax-design questions AI raises, including whether existing bases can reach AI-driven rents.

Microsimulation of AI shocks. Closest in method are two companion studies on European systems. Doorley et al. (2026) run scenario-based AI employment and wage shocks through the Irish SWITCH model, and Ahmadi (2026) runs occupation-level exposure-based displacement, wage and capital shocks through PolicyEngine UK, treating the incidence of displacement as an explicit scenario axis. This paper is the United States member of that family, but it differs in the shock design: rather than displacement incidence, it varies the factor composition and wage-distribution shape of income growth, matching the Budget Lab’s calibration to enable the cell-level reconciliation of Section 6. Occupation-level exposure measures (Felten et al., 2021; Eloundou et al., 2024) are the natural next refinement of the wage channel here, as Section 7 discusses; both this paper’s repository and the Budget Lab’s roadmap point the same direction.

Factor shares and automation. The premise that technology can move income between labor and capital at scale is not hypothetical: the global decline of the labor share is documented by Karabarbounis and Neiman (2014), and Acemoglu and Restrepo (2018) provide the task-model mechanics by which automation reallocates value added from labor to capital. What

microsimulation adds to that literature is the statutory detail: a dollar moved from wages to capital income changes its withholding, its marginal rate schedule, its payroll treatment, its state treatment, and the transfer payments of the household it leaves — none of which aggregate models carry. The scenarios here take the CBO baseline as the no-AI counterfactual, using the same February 2026 projections the Budget Lab anchors to ([Congressional Budget Office, 2026](#)).

3 Model, data, and experiment designs

3.1 Model and data

All results come from PolicyEngine-US version 1.764.6, an open-source static microsimulation of United States federal and state tax and benefit law, executed through the `policyengine.py` runtime (version 5.0.1), which certifies the pairing of model version and dataset build. The dataset is `populace-us` build `cae8640` (28 July 2026), an openly published enhanced-CPS microdata file for 2024, uprated by the model to the analysis years. The mechanism sweep of Section 4.1 is run at 2026; the calibrated scenarios are run at 2030, matching the Budget Lab’s horizon. In the 2030 baseline, SPM poverty is 12.13% and the household net-income Gini is 0.4625 (household-level, unequivalized; Section 7 discusses comparability).

The model covers federal individual income tax (including capital gains rates, the net investment income tax and the alternative minimum tax), payroll taxes on both employee and employer sides plus self-employment tax, refundable credits (EITC, refundable CTC and others), the major means-tested transfer programs (SNAP, SSI, TANF, WIC, school meals, housing subsidies, broadband subsidies), and the tax and benefit codes of all fifty states and DC. Health programs (Medicaid, CHIP, ACA subsidies) are valued in the model but excluded from net-income accounting in these runs by the model’s default simulation switch, so no result here contains a health-program response in either direction. Everything is static: no labor supply, behavioral, avoidance or enforcement responses, matching the Budget Lab’s individual microsimulation.

Accounting. Net government revenue is defined as taxes minus transfers: household tax before refundable credits (federal and state, including employee-side payroll) plus employer payroll tax, minus refundable credits and benefit outlays. An identity check — the change in market income minus the change in household net income must equal the household-side revenue change — closes to \$0.000B in every scenario. Where the paper compares against the Budget Lab (Section 6), it restates results on their narrower basis: federal individual income tax net of refundable credits, plus payroll.

3.2 Design I: the mechanism sweep

The sweep isolates composition. In 10-point steps from 0% to 100%, it reduces every person’s positive employment and self-employment income by the step share and redistributes the same weighted total to positive capital income in proportion to existing holdings, holding modelled market income fixed. The 2026 baseline contains \$12.4T of positive labor income and \$2.7T of positive capital income, with 51% of households receiving any positive capital income — so the sweep concentrates income by construction, and the question is what current law does about it.

The sweep’s limitation is deliberate: a pure reallocation holds output fixed, while forecasters expect AI to raise output *and* tilt its composition. The sweep therefore answers “what does composition alone do to current law?”, and the calibrated scenarios answer “what happens on the paths forecasters actually expect?”. The two connect through a forecast-equivalence mapping:

the share of baseline labor income that the tilt removes, $k = 1 - (1 + g_L)/(1 + g_Y)$, places each scenario on the sweep axis at 0.9% (Slow), 3.1% (Moderate) and 7.6% (Rapid) — the sweep’s forecast-relevant range is its first decade, and the rest is out-of-sample stress test.

3.3 Design II: forecaster-calibrated scenarios

The scenarios adopt the Budget Lab’s calibration of Karger et al. (2026) exactly. Each scenario is defined by an annual real GDP growth rate under AI (Slow 2.0%, Moderate 2.6%, Rapid 3.3%) and a 2030 labor share of factor income (55.0%, 53.8%, 51.3%). Growth is expressed in excess of the CBO baseline path: the five-year output shock is $g_Y = (1 + r_{ai})^5/G_{CBO} - 1$, and the factor shocks follow from the target shares,

$$g_K = \frac{\theta_K^1(1 + g_Y) - \theta_K^0}{\theta_K^0}, \quad g_L = \frac{\theta_L^1(1 + g_Y) - \theta_L^0}{\theta_L^0}. \quad (1)$$

Two inputs their report uses but does not print — the cumulative CBO baseline factor $G_{CBO} = 1.0976$ and the pre-shock labor share $\theta_L^0 = 0.5550$ — are recovered by inverting their published formulas against the nine derived values in their Table 1; unit tests reproduce all nine to within 0.005 pp and re-derive the labor share by grid search so the constants cannot drift.

The wage-inequality channel is their stylization $\lambda = 1 \pm g_Y$: a spread transform on log labor income, $\ln Y_i^1 = \mu^1 + \lambda(\ln Y_i^0 - \mu^0)$, equivalently $Y_i^1 = C Y_i^{0\lambda}$ with C in closed form, with the aggregate pinned to $(1 + g_L)$ times baseline. Compressive ($\lambda < 1$) narrows the wage distribution around the same aggregate; expansive ($\lambda > 1$) widens it; proportional scales every wage equally. The transform is applied at the person level because the payroll result turns on a per-person threshold, the Social Security taxable maximum (\$214,200 in 2030; at baseline 15.8% of wage income and 5.7% of workers sit above it).

The capital shock scales nine capital income variables — long- and short-term capital gains, taxable interest, qualified and non-qualified dividends, rental income, partnership and S-corporation income, taxable pension income and taxable IRA distributions — in proportion to existing holdings, after routing pass-through income through the Budget Lab’s labor/capital split rule. Appendix B traces every shocked variable into IRS gross income and household market income, confirming the shock reaches the tax base rather than being recomputed away; the labor aggregate lands on its target to within relative error 2e-16. Each scenario also has a shares-fixed counterfactual growing both factors at g_Y , the comparison behind the Budget Lab’s “roughly twice as large” finding.

Two further axes vary definitional assumptions their report flags. A *realization rate* routes a fraction of the incremental capital flow outside the taxable base (their assumption corresponds to 1.0); and a *capital-scope* variant excludes taxable pension income and IRA distributions from the capital set, since those are taxed at ordinary rates and their inclusion works against the “capital is taxed lightly” mechanism.

4 Results

4.1 The mechanism at full range

Figure 1 runs the composition mechanism from 0% to 100%. Net government revenue falls almost monotonically as labor income becomes capital income, bottoming at -880B at a 90% shift: the income tax initially loses progressivity revenue as wages leave the schedule, payroll revenue falls




Figure 1: Revenue under a 0–100% labor-to-capital shift, 2026. Lines are changes against baseline; benefit outlays are sign-flipped so a rising line is a fiscal cost. The shaded band marks the forecast-equivalent positions of the 2030 scenarios.

roughly linearly with the wage base, state codes track the federal pattern, and benefit outlays rise as wage-dependent households qualify for means-tested support. The distributional counterpart (Figure 2) is severe by construction — only 51% of households hold any positive capital income, so the net-income Gini rises from 0.470 to 0.768 and SPM poverty from 13.2% to 61.2% at the full shift, with the top decile’s mean net income rising +58% while decile 5 falls -48.5%. The forecast-equivalent range — the shaded band at 0.9–7.6% — says how much of this axis the 2030 forecasts actually reach: the first decade. The sweep’s value is the mechanism it exposes, not the tail.

4.2 Calibrated scenarios

Table 1 reports all twelve scenario cells at 2030 against the same-year current-law baseline; Figure 3 plots the Rapid pattern. Two published Budget Lab mechanisms reproduce on independent machinery. Within each scenario, moving from compressive to expansive raises income-tax revenue through progressivity and lowers payroll revenue as earnings cross the Social Security taxable maximum. And the tilt is expensive: Table 2 shows each scenario keeping a fraction of its shares-fixed revenue — 36% in Rapid (+211B against +581B).

The payroll-cap mechanism runs in both directions, which is worth stating because it is easy to get backwards: under Rapid / compressive, aggregate labor income *falls* -0.94% and payroll receipts still *rise* +34B, because compression moves earnings from above the taxable maximum to below it; under expansive the same cap turns wage growth into a payroll loss of -70B. The wage-spread parameter also has a direct household reading: each variant has a crossover wage below which workers lose (expansive, \$168,205) or gain (compressive, \$118,287), so under Rapid / expansive the great majority of workers sit on the losing side of a scenario that adds +244B to net revenue.




Figure 2: Inequality and poverty under the same sweep.

4.3 The transfer system

The Revenue column of Table 1 nets transfers out; pulling them apart shows the safety net responding exactly as designed (Figure 4). Across the Rapid wage-spread variants — identical aggregate labor income, only its distribution varies — gross tax collections swing \$122B, and automatic transfer responses absorb \$57B of that swing, 47%. Non-credit transfers alone swing \$42B: -20B under compressive, as rising low-end wages float households off means-tested programs, to +22B under expansive as falling ones push them on. SNAP carries \$19B of the swing (-9B to +10B) and SSI \$3B; Moderate shows the same shape (-11B to +8B). Refundable credits contribute the remaining \$15B — a response the Budget Lab’s accounting does capture through its tax-credit outlays line. The \$42B of program spending is the piece outside their model’s coverage, and it is the outlay half of the 3.26pp poverty span: the same variant uncertainty their report scores as a revenue question arrives at households as a benefits question.

4.4 States

State codes are the other channel a federal model does not carry. Under Rapid / proportional they collect +36B, and the state take doubles across the wage-spread variants (+24B to +48B) as progressive state schedules concentrate collections at the top the way federal brackets do. On the Rapid / expansive cell used for the reconciliation below, state tax net of refundable state credits adds +47B — roughly a sixth on top of the bridged federal take.

The geography is concentrated (Figure 5): the top five states collect 56% of the national state take (CA +10.0, NY +5.8, NJ +1.6, MA +1.3, MD +1.2, in \$B). 7 states collect exactly nothing from the capital shock — Florida, Nevada, New Hampshire, South Dakota, Tennessee, Texas and Wyoming — and Washington is the instructive exception: it has no income tax either, but its capital-gains excise tax collects +1.1B, more than the seven zero-collecting states combined. A state that taxes capital gains participates in an AI capital boom; a state that taxes only wages does not participate in an economy where wages are the shrinking share. State-funded benefit responses are negligible in these runs (the model’s state-benefit registry covers state SSI supplements and child-care subsidies), so the automatic-stabilizer response above is almost

Scenario	Revenue	Income tax	Payroll	State	Poverty	Δ pp	Net Gini
Slow / shares fixed	+47	+29	+10	+6	11.99%	-0.14	0.4627
Slow / compressive	+3	+3	-1	+1	12.01%	-0.12	0.4622
Slow / proportional	+6	+11	-5	+2	12.17%	+0.04	0.4632
Slow / expansive	+10	+19	-10	+3	12.33%	+0.20	0.4643
Moderate / shares fixed	+289	+178	+61	+35	11.46%	-0.67	0.4640
Moderate / compressive	+128	+65	+33	+15	11.18%	-0.95	0.4596
Moderate / proportional	+144	+114	+7	+21	12.00%	-0.13	0.4656
Moderate / expansive	+163	+164	-20	+28	12.84%	+0.71	0.4718
Rapid / shares fixed	+581	+358	+121	+71	11.04%	-1.09	0.4654
Rapid / compressive	+179	+98	+34	+24	10.59%	-1.54	0.4581
Rapid / proportional	+211	+193	-16	+36	12.24%	+0.11	0.4699
Rapid / expansive	+244	+294	-70	+48	13.85%	+1.72	0.4820

Table 1: Scenario grid, 2030, \$B change against the current-law baseline (SPM poverty 12.13%, net Gini 0.4625). “Revenue” is net government revenue (federal and state taxes minus refundable credits and benefits); “Income tax” is federal income tax gross of refundable credits; “State” is state tax before refundable state credits.

Scenario	Shares fixed	As forecast	Share kept
Slow	+47	+6	14%
Moderate	+289	+144	50%
Rapid	+581	+211	36%

Table 2: What the tilt toward capital costs: net revenue with factor shares fixed versus as forecast (proportional wage spread), \$B, 2030.

entirely federal.

5 The two assumptions that dominate

5.1 Realization

The Budget Lab’s capital shock assumes AI leaves the realized share of capital income unchanged, and they flag the assumption as a limitation. It carries real weight: unrealized gains, retained earnings and accruals inside retirement accounts never reach the individual tax base. Holding Rapid / proportional fixed and varying only how much of the incremental capital flow lands on a tax return (Figure 6, Table 5 in the appendix), revenue runs from -61B at a 0% realized share to +211B at their 100% assumption — linear to within a percent, with breakeven at a 22.5% realized share. Below that, the household side of the fiscal system loses money on Rapid AI growth: labor income still falls, and not enough taxable capital income reaches returns to offset it.

One counterweight sits outside the household sector: corporate tax is levied upstream of household realization, so the Budget Lab’s corporate wedge would still collect about +84B on the same shock regardless of the realized share — enough to keep the all-in federal total positive, about +23B, even at the 0% floor. Low realization shifts the fiscal burden of an AI boom onto




Figure 3: Net revenue (left) and SPM poverty change (right) by scenario and wage-inequality variant, 2030.

the corporate side; it does not eliminate it.

Poverty is flat across the entire sweep, between 12.24% and 12.27%. Whether AI’s capital gains are taxable moves the federal balance sheet by more than a quarter-trillion dollars and moves poor households essentially not at all: the transfer system responds to what happens to wages, not to how much of the capital gain the tax system reaches.

5.2 What counts as capital

Taxable pension and IRA distributions are taxed at ordinary rates. Counting them in the capital set — which the Budget Lab does, and the headline results follow — raises the revenue estimate and softens the “capital is taxed lightly” mechanism. Excluding them (Table 6 in the appendix) moves Rapid from +211B to +106B and flips Slow’s sign (+6B to -4B). Together with realization, these two definitional choices span more of the revenue answer than the choice between Slow, Moderate and Rapid.

5.3 Stability across data builds

The results were first computed on the prior **populace-us** build (build o, 22 July 2026) and rerun on build p (28 July, with revised capital gains calibration) under the identical pinned model version — a data-only change. Every scenario revenue figure moved by less than 3%, every poverty change by less than 0.11pp, and no conclusion changed. The recalibration’s visible effect is distributional: the baseline top-1% net income share rises from 8.77% to 9.00%, and New York’s state take under Rapid rises from \$3.8B to \$5.8B as more of the gains land where they are taxed. An average-tax-rate cross-check also holds: Rapid / proportional turns a +760B market income increase into +211B of receipts, 27.8%, inside the 21–34% band the Budget Lab reports on gross factor income.




Figure 4: Gross tax collections, automatic transfer response, and net revenue across the Rapid wage-inequality variants, 2030. Transfers (refundable credits plus benefits) rise exactly when workers lose, absorbing 47% of the gross-tax swing across the variants.

Rapid / expansive	IIT (net)	Payroll	Corporate wedge	Total
This paper	+287	-70	+84	+301
Budget Lab	+195	-63	+84	+216

Table 3: The bridged headline cell (Rapid / expansive, their central published scenario), \$B, 2030. “IIT (net)” is individual income tax net of refundable credits; the corporate wedge is theirs in both rows.

6 Reconciliation with the Budget Lab

The two models measure different totals: theirs is federal individual income tax plus payroll plus a corporate wedge, with no states and no non-credit transfers; this paper’s household total spans federal and state and nets out the whole transfer system. Setting the headline numbers side by side without a bridge would be wrong, so this section builds one, using only this paper’s run output and objects they publish. Their repository commits the complete published result grid — revenue by instrument for all eighteen cells — which makes the comparison possible at the cell level rather than against two text-stated totals.

Basis. Their microsimulation total is individual income tax plus payroll minus tax-credit outlays; the matching quantity here is federal income tax net of refundable credits, plus payroll (both halves and self-employment). Their corporate wedge reduces to the capital growth rate times CBO’s 2030 corporate anchor of \$486B — in their formula $\Delta R_{CIT} = X \cdot (R_{CIT}/Y_K^0)$ with $X = Y_K^0 g$, the capital base cancels — and all eighteen committed values equal $g \times \$486B$ to within \$0.15B. Adding their wedge to this paper’s federal microsimulation total yields a like-for-like “bridged” headline. The wedge is not PolicyEngine modelling corporate tax, and it inherits every constant-rate assumption they list.

Where the models agree. Payroll responses agree closely in every Rapid cell (Table 4):




Figure 5: Top ten states by net state revenue change, Rapid / proportional, 2030.

Rapid variant	IIT net of credits		Payroll	
	Ours	Theirs	Ours	Theirs
Compressive	+104	+11	+34	+37
Proportional	+193	+98	-16	-11
Expansive	+287	+195	-70	-63

Table 4: Cell-level comparison on the matched federal basis, Rapid variants, \$B, 2030.

ours run +34/-16/-70 against theirs +37/-11/-63 across compressive / proportional / expansive — two independent models putting similar wage distributions against the same taxable maximum, including the sign reversal under compression. The tilt ratios also line up once stated on the same basis: on their federal-plus-corporate accounting, shares-fixed revenue exceeds as-forecast revenue by $3.13\times$ / $1.65\times$ / $2.00\times$ here against $3.62\times$ / $1.71\times$ / $2.11\times$ in their grid for Slow / Moderate / Rapid — agreement on the paper’s central mechanism, with the same ordering. (The household-total ratios here are larger — $2.75\times$ for Rapid — because states and transfers amplify what the tilt costs; that is a coverage statement, not a disagreement.) Their most extreme published comparison also reproduces directionally: for Slow with compressive inequality, their grid puts revenue 84% below the shares-and-distribution-unchanged counterfactual (their text states 82%); this paper gets 93%.

Where they diverge, and why. The bridged headline is +301B against their +216B — $1.40\times$, or 4.6% against 3.3% of CBO 2030 revenue. The gap concentrates in the income-tax response to the capital shock: on the proportional cell, where the wage-spread channel is neutral and the income-tax response is almost entirely the capital shock, this paper’s response is $1.97\times$ theirs ($1.47\times$ on the expansive bridge cell). Three model differences point the same direction. First, the realized-capital base: their methodology’s remark that the effective corporate parameter “falls out to roughly one” implies a baseline realized taxable capital base near \$4.6T, against \$5.3T here, roughly 14% larger. Second, allocation: they distribute the incremental




Figure 6: Net revenue under Rapid / proportional as the realized share of the incremental capital flow varies. The dashed line adds the Budget Lab’s realization-independent corporate wedge.

capital flow by SCF-imputed total wealth; this paper distributes by existing realized capital income, which they note is the more top-concentrated choice — and top-concentrated dollars face higher marginal rates. Third, retirement routing: taxable pension and IRA distributions here scale with the capital shock and are taxed at ordinary rates. The capital-scope sensitivity of Section 5.2 shows that last channel alone is worth about \$105B in Rapid.

7 Discussion

7.1 What the results say

Three readings of the results bear directly on policy questions the AI-fiscal conversation is starting to ask. First, the fiscal and household questions decouple. Capital-side tax design (the realization margin) moves the budget by hundreds of billions while leaving poverty untouched; the wage-side distribution moves poverty by 3.26pp while net revenue stays positive throughout. A revenue estimate, however good, is not a household forecast. Second, the automatic stabilizers are already in the scenario arithmetic: transfers absorb 47% of the gross-tax swing across the wage-spread variants, which means deficit projections that stop at the tax code overstate how much the wage-distribution assumption matters fiscally — and understate the program-side costs when workers lose. Third, the state layer is not a rounding error: +36B under proportional Rapid growth, doubling across wage-spread variants, concentrated in a handful of states, and available only to states whose codes reach capital income at all.

7.2 Limitations

No corporate sector. PolicyEngine-US models households; the corporate wedge in Section 6 is borrowed from the Budget Lab, not modelled.

Static. No labor supply, saving, avoidance or enforcement responses on either side of the comparison.

Apportionment is narrower than theirs. The incremental capital flow follows existing realized capital income; their SCF-wealth apportionment reaches households that hold wealth but realize nothing. By their own account the narrower base is more top-concentrated, so this choice concentrates the shock more than theirs; the microdata would support an asset-based variant (`net_worth` is populated for 99.7% of households), which is identified future work.

Pass-through split is coarser than theirs. Their active/passive-aware rule is approximated by a single rule because the microdata carries no active/passive flag, under-assigning the passive portion to capital.

National shares versus modelled shares. Scenario calibration uses national factor shares ($\theta_L^0 = 0.5550$); observed taxable microdata flows are more labor-weighted (\$15.9T labor against \$5.3T capital), because survey and tax data miss unrealized gains, retained earnings and imputed returns. Growth rates from national accounts are applied to observed taxable flows — the same choice the Budget Lab makes with the PUF — and both baseline shares are reported so the gap is visible rather than assumed away.

Anchored poverty, unequivalized Gini. SPM thresholds do not respond to the simulated distribution, so a fully relative poverty line would rise with five years of above-baseline growth and show larger increases. The Gini is household-level and unequivalized: changes are meaningful, levels are not comparable to published equivalized series. Poverty levels sit above the Census published rate — a known calibration gap — so every poverty figure is a change against a same-year baseline on identical thresholds.

No health-program response. Medicaid, CHIP and ACA subsidies are excluded from net-income accounting in these runs by the model’s default switch; the transfer responses reported here would not shrink, and could grow, with that channel on — but this paper does not measure it.

7.3 Next steps

The wage-inequality parameter λ is a stylization with, as the Budget Lab says plainly, little evidence behind its magnitude. The natural replacement is measured occupation-level exposure (Felten et al., 2021; Eloundou et al., 2024): distributing the labor shock by exposure scores mapped to CPS occupations would make the implied λ an output rather than an input — the design of the UK companion paper (Ahmadi, 2026), and the “S1” expansion the Budget Lab’s methodology names as planned future work. Asset-based capital apportionment and a policy-response layer (recycling the net revenue gain into transfers to test whether the receipts can hold the bottom of the distribution harmless) are the other two identified extensions.

7.4 Conclusion

Whether AI pays for itself fiscally depends on where its income lands — between factors, across the wage distribution, and on or off the individual tax base. Running forecaster-calibrated scenarios through the full United States tax and transfer system, the tilt toward capital keeps 36% of the revenue that balanced growth would deliver and erases growth’s poverty gain; the wage-distribution assumption nobody can yet pin down flips the household verdict while the revenue verdict stands; and the channels a tax-only model omits — automatic transfers and state codes — carry a material share of the response. Where this model and the Budget Lab’s overlap,

they agree closely; where they diverge, the anatomy is identifiable and mostly definitional. All of it is open: anyone can rerun these scenarios, or score a different future against them.

References

- Acemoglu, D. (2025). The simple macroeconomics of AI. *Economic Policy*, 40(121), 13–58.
- Acemoglu, D. and Restrepo, P. (2018). The race between man and machine: Implications of technology for growth, factor shares, and employment. *American Economic Review*, 108(6), 1488–1542.
- Ahmadi, V. (2026). Who bears the AI shock? The distributional and fiscal incidence in the UK. SSRN Working Paper No. 7174479. https://papers.ssrn.com/sol3/papers.cfm?abstract_id=7174479.
- Bastani, S. and Waldenström, D. (2024). AI, automation and taxation. CESifo Working Paper No. 11084 / IZA Policy Paper No. 212, Munich and Bonn.
- Congressional Budget Office (2026). *The Budget and Economic Outlook: 2026 to 2036*. Washington, DC, February 2026. <https://www.cbo.gov/publication/61882>.
- Doorley, K., O’Connor, S., O’Shea, R. and Tuda, D. (2026). Artificial intelligence and income inequality in Ireland. Jointly Published Report No. 16, Economic and Social Research Institute and Department of Finance, Dublin. <https://doi.org/10.26504/jr16>.
- Eloundou, T., Manning, S., Mishkin, P. and Rock, D. (2024). GPTs are GPTs: Labor market impact potential of LLMs. *Science*, 384(6702), 1306–1308.
- Felten, E., Raj, M. and Seamans, R. (2021). Occupational, industry, and geographic exposure to artificial intelligence: A novel dataset and its potential uses. *Strategic Management Journal*, 42(12), 2195–2217.
- Iselin, J. and Nunn, R. (2026). How potential AI futures would play out in the current tax system. The Budget Lab at Yale, New Haven, CT, 20 July 2026. <https://budgetlab.yale.edu/research/how-potential-ai-futures-would-play-out-current-tax-system>.
- Karabarbounis, L. and Neiman, B. (2014). The global decline of the labor share. *The Quarterly Journal of Economics*, 129(1), 61–103.
- Karger, E., Kuusela, O., Abaluck, J., Bryan, K. A., Halperin, B., Jones, T. R., Murphy, C., Trammell, P., Reynolds, M., Mayland, D., Viswanathan, R., Mittal, A., Ceppas de Castro, R., Rosenberg, J. and Tetlock, P. (2026). Forecasting the economic effects of AI. NBER Working Paper No. 35046, National Bureau of Economic Research, Cambridge, MA.

A Additional tables

Realized share	Revenue
0%	-61
25%	+7
50%	+74
75%	+142
100% (their assumption)	+211

Table 5: Net revenue under Rapid / proportional by realized share of the incremental capital flow, \$B, 2030.

Scenario	Full capital set	Excl. retirement	Difference
Slow	+6	-4	-10
Moderate	+144	+98	-46
Rapid	+211	+106	-105

Table 6: Capital-scope sensitivity: net revenue with and without taxable pension and IRA distributions in the capital set (proportional variants), \$B, 2030.

State	Net revenue change
CA	+10.01
NY	+5.82
NJ	+1.58
MA	+1.33
MD	+1.16
VA	+1.15
WA	+1.13
MN	+1.02
IL	+0.91
OR	+0.89
7 states collect exactly \$0: FL, NV, NH, SD, TN, TX, WY.	

Table 7: Top ten states by net state revenue change (state tax before refundable credits minus refundable state credits), Rapid / proportional, \$B, 2030.

Shift	Revenue	Income tax	Poverty	Net Gini
0%	+0	+0	13.2%	0.470
10%	-135	+44	15.0%	0.490
20%	-255	+101	16.9%	0.513
30%	-370	+176	20.1%	0.536
40%	-486	+265	23.9%	0.561
50%	-592	+371	29.3%	0.588
60%	-688	+493	35.4%	0.617
70%	-773	+638	43.1%	0.647
80%	-847	+807	50.7%	0.681
90%	-880	+1000	57.6%	0.719
100%	-790	+1235	61.2%	0.768

Table 8: The mechanism sweep, 2026: net revenue, federal income tax, SPM poverty and net Gini by share of labor income shifted to capital.

B Shock routing into the tax base

Several shocked variables are aggregates rather than raw inputs, so the implementation verifies that overriding them moves taxable income rather than being recomputed away. Traced against the model's gross-income and market-income source registries (`gov.irs.gross_income.sources` and `gov.household.market_income_sources` in policyengine-us 1.764.6):

Shocked variable	Route into IRS gross income	Route into market income
employment income	via employment income, net of pre-tax contributions	listed source
self-employment income	listed source	listed source
long/short-term capital gains	capital gains	capital gains
qualified/non-qualified dividends	ordinary dividend income	dividend income
taxable interest	listed source	interest income
rental income	listed source	listed source
partnership and S-corp income	listed source	listed source
taxable pension income	listed source	pension income
taxable IRA distributions	taxable retirement distributions	retirement distributions

Table 9: Every shocked variable reaches both the federal tax base and household market income. Qualified and non-qualified dividends are the two components of ordinary dividends; short- and long-term gains the two components of capital gains; tax-exempt interest is deliberately outside the shocked set.

Because every shocked variable reaches household market income, the net-income identity closes; the maximum residual across scenarios is \$0.000B, and the labor aggregate lands on its target to within relative error 2e-16.

C Reproducibility

Everything in this paper runs from open source against openly published microdata:

- **Model:** policyengine-us 1.764.6 via policyengine.py 5.0.1 (the runtime certifies the model–data pairing).
- **Data:** populace-us certified build populace-us-2024-buildp-sparse-rmloss100-cae8640-20260728T0
- **Code:** github.com/PolicyEngine/ai-inequality — scenario engine in `analysis/ai_scenarios.py`, runner in `analysis/compute_ai_scenarios.py`, sweep in `analysis/compute_shift_sweep.py`, reconciliation in `analysis/reconcile_budget_lab.py`. Every number and table in this paper is emitted from the run outputs by `analysis/emit_paper_values.py`; figures by `analysis/paper_figures.py`.
- **Their side:** the Budget Lab’s AI-Fiscal repository (MIT license) commits the complete published result grid (`results/aggregates/ai_fiscal_publishable_2030_latest.xlsx`); the reconciliation reads it directly. Their PUF-based input microdata cannot be redistributed, so their pipeline is readable but not rerunnable by third parties; this paper’s is both.
- **Interactive companion:** <https://policyengine.org/us/ai-inequality/income-shift> (scenario explorer) and <https://policyengine.org/us/ai-inequality/budget-lab> (memo-length summary).

Calibration unit tests reproduce all nine of the Budget Lab’s published derived scenario values to within 0.005 pp from the two recovered constants ($G_{\text{CBO}} = 1.0976$, $\theta_L^0 = 0.5550$) and re-derive the labor share by grid search. The scenario runner checkpoints per-scenario results and is resumable; run logs and build-o archives are kept under `analysis/outputs/`.